Journal of Behavioral and Experimental Economics 112 (2024) 102239

Contents lists available at ScienceDirect

Journal of Behavioral and Experimental Economics

journal homepage: www.elsevier.com/locate/jbee

Do people trust humans more than ChatGPT?

Joy Buchanan [a,*], William Hickman [b]

[a] Brock School of Business, Samford University, Birmingham, AL, United States of America
[b] The Interdisciplinary Center for Economic Science, George Mason University, Arlington, VA, United States of America

ARTICLE INFO

JEL Codes:
O33
C91
D8

Keywords:
Artificial intelligence
Machine learning
Trust
Belief
Experiments

ABSTRACT

We explore whether people trust the accuracy of statements produced by large language models (LLMs) versus those written by humans. While LLMs have showcased impressive capabilities in generating text, concerns have been raised regarding the potential for misinformation, bias, or false responses. In this experiment, participants rate the accuracy of statements under different information conditions. Participants who are not explicitly informed of authorship tend to trust statements they believe are human-written more than those attributed to ChatGPT. However, when informed about authorship, participants show equal skepticism towards both human and AI writers. Informed participants are, overall, more likely to choose costly fact-checking. These outcomes suggest that trust in AI-generated content is context-dependent.

1. Introduction

The capabilities of Large Language Models (LLMs) have significantly improved in recent years, transforming natural language processing and giving LLMs the ability to generate human-like text with remarkable fluency. In particular, OpenAI's ChatGPT has demonstrated the ability to generate coherent and contextually relevant responses to various prompts. While LLMs have displayed impressive capabilities, concerns have been raised regarding the potential for misinformation or bias in their responses. Indeed, ChatGPT is known to produce inaccurate results sometimes. Understanding the level of trust humans place in LLMs is crucial, as these models are increasingly being deployed in various real-world applications, including content generation, automated customer support, and information retrieval.

LLM-based tools will be adopted where they are cost-effective because of demonstrated successes in workplace settings (Brynjolfsson, Li, & Raymond, 2023; Fumagalli, Rezaei, & Salomons, 2022). So, it is increasingly likely that human readers will encounter AI-written text.

We use an incentivized survey experiment to investigate whether humans trust the factual accuracy of LLM output. Our study is the first controlled experiment of this type that is incentivized and gives the human participants the option of doing a costly fact-check, which mirrors the options available to workers and managers in the workplace. Participants rate the accuracy of statements written by humans or ChatGPT. We vary whether each statement was written by a human or ChatGPT and whether we explicitly inform participants about authorship.

Our findings paint a complex portrait of trust. In this setting, trust is defined as believing in the credibility or accuracy of a statement. If participants are not informed about authorship but first have to guess whether a paragraph was written by AI or by a human, then the participants who guess the paragraph was written by a human display higher trust in the statement's factual accuracy. However, in the treatments with explicit information about authorship, participants are equally suspicious that the AI or human-written paragraph contains an error and needs to be fact-checked. We collect information about age and whether participants have had exposure to ChatGPT prior to the experiment. Neither of those characteristics has a significant impact on trust.

Our results indicate that trust in content believed to be AI-generated is context-dependent, especially now that AI writing can closely mimic human style and content. We add to a growing literature finding that informed individuals are not inherently biased against the accuracy of AI outputs (Epstein et al., 2023). Our emphasis is on trust in AI to perform an objective task; however, there might be different results in an experiment that measure attitudes toward AI performing subjective tasks (Castelo, Bos, & Lehmann, 2019). We are using factual accuracy in writing as this objective task because the recent increase in the use of LLMs to write original text makes this a relevant task to assess.

This research will provide insights for researchers, developers, and

* Corresponding author.
E-mail address: jbuchan1@samford.edu (J. Buchanan).

https://doi.org/10.1016/j.socec.2024.102239
Received 22 November 2023; Received in revised form 18 March 2024; Accepted 21 May 2024
Available online 31 May 2024
2214-8043/© 2024 Elsevier Inc. All rights are reserved, including those for text and data mining, AI training, and similar technologies.

policymakers, paving the way for a more informed integration of LLMs into our increasingly AI-driven world. We also present a novel design that can be extended and modified to study trust in AI-based tools.

The remainder of this paper is structured as follows. Section 2 discusses the literature related to our research. Section 3 explains our experimental design and procedures, while Section 4 details our hypotheses. Section 5 presents our results, and Section 6 concludes.

2. Related literature

LLMs can "hallucinate" falsehoods that misinform human readers. In a survey article, Ji et al. (2023) state that "deep learning based generation is prone to hallucinate unintended text, which degrades the system performance and fails to meet user expectations in many real-world scenarios." Zhang et al. (2023) find that LLMs cannot fully self-correct or recognize mistakes. Buchanan, Hill, and Shapoval (2024) demonstrated that ChatGPT produces fake citations for papers that do not exist, with both GPT-3.5 and the more advanced GPT-4. Open AI researchers concluded after running image processing tests that "Given the model's imperfect performance in this domain and the risks associated with inaccuracies, we do not consider the current version of GPT-4 V to be fit for performing any medical function or substituting professional medical advice, diagnosis, or treatment, or judgment" (OpenAI, 2023). Regardless, it is likely that people will use LLMs for many tasks, including medical and policy advice. Measuring trust in LLMs is needed in order to gauge the impact they will have on real-world outcomes.

While Gillespie et al. (2023) and others have conducted surveys about trust in AI, we are among the first to do an incentivized study to test whether people trust AI generated writing. Scientists should not rely completely on self-reported data for this issue (Smith, 1994). Humans might think they are supposed to say humans are more trustworthy than AI in a survey. However, if fact checking is costly, then these same people might rely on AI-generated text in practice without screening it for inaccuracies.

In an unincentivized study, Spitale, Biller-Andorno, and Germani (2023) recruit human participants to evaluate tweets written by humans versus tweets written by AI. They conclude that, "In comparison with humans, [LLMs] can produce accurate information that is easier to understand, but it can also produce more compelling disinformation. We also show that humans cannot distinguish between tweets generated by GPT-3 and written by real Twitter users." Köbis and Mossink (2021) study whether respondents have a subjective preference for art created by humans versus poetry written by AI, with mixed results. Babin and Chauhan (2023) find that people prefer advice from humans in social dilemmas and have a higher willingness to pay for advice from humans.

To date, there has been more experimental work on lying than on unintentional AI mistakes. For example, Chen and Houser (2017) examine whether human participants trust messages written by human counterparts in a trust game. Serra-Garcia and Gneezy (2021) ask people to detect messages with lies in an incentivized study. They find that participants are overconfident in their own ability to detect lies (from audio visual sources) which has implications for what gets amplified on social media.

Pennycook et al. (2021) use a prompt that causes people to pause and think about accuracy before sharing a story on social media. They find that "subtly shifting attention to accuracy increases the quality of news that people subsequently share."[1] Our results might be interpreted as a replication of Pennycook et al. (2021) because when we draw attention to authorship, our participants become suspicious and more inclined to do a costly fact-check.

Chugunova and Sele (2022) and March (2021) provide overviews of studies on human interactions with autonomous agents. Humans treat AI counterparts differently than human strategic partners. It is evident that humans apply different cognitive processes when they know they are facing a computer.[2] Chugunova and Sele (2022) state that "while humans seem willing to accept automated agents in areas considered more objective or analytical, they seem reluctant to do so in areas considered social or moral." We replicate that to some extent when we test human perception of an AI doing a fairly objective task which is to write a factual statement about an uncontroversial historical topic. Our participants are about as suspicious of ChatGPT as human writers, when they are informed about authorship. However, we find that context affects the willingness to trust AI. This finding is consistent with the finding that humans are willing to incorporate the advice of computers (Leib et al., 2023). Yet, this willingness varies based on the framing of the situation and the perception of autonomy. Sunstein and Reisch (2023) also find that framing affects the human perception of algorithms.[3]

Because LLMs are new, we seek to provide novel evidence on how humans perceive natural language that they believe was entirely written by a computer.

3. Experimental design and procedures

Participants answered questions after reading a 5-sentence paragraph. Two types of paragraphs were used: human-written and AI-generated. The human-written text came from an entry in

Encyclopedia Britannica that discusses the history of statistics and probability (Porter, 2023). The AI-written text was generated by prompting ChatGPT to replicate the style, sentence count, and word count of the human-written paragraph without altering factual content.[4] Fig. 1 shows Paragraph A, written by a human, as it was presented to participants who were explicitly informed of human authorship. Appendix B shows a copy of the survey, as we presented it to participants who were explicitly informed of AI authorship (for Paragraph A).

We randomly assigned participants to one of the treatments listed in Table 1.

Participants indicated whether they thought the paragraph was written by a human or AI, with an incentive for being correct. If the experimenter informed the participants about authorship, then this question served as an attention check, and otherwise it was an incentivized belief elicitation.

We did not give uninformed participants any additional information about the probability their paragraph was human or AI-written.

Participants could earn money for correctly guessing whether the paragraph contained any factual errors. They could pick from three responses:

1. There are not any factual errors in this paragraph
2. I would like to purchase a fact-check (cost: $0.20)
3. There are factual errors in this paragraph

Correct answers without a fact-check earned a $0.50 bonus. Opting for a fact-check guaranteed a $0.30 bonus ($0.50 - $0.20).

[1] This result was replicated by Athey et al., (2023), who ran phone-based tutorials in Kenya about not impulsively sharing stories on social media that might be false.

[2] Buchanan (2022) demonstrates the importance of subjective preferences in determining how humans engage with computers and decide to invest in training.

[3] Smith and Wilson, (2019) develop a model for understanding how context influences behavior.

[4] To ensure the AI-written text was similar to the human-written text, we prompted ChatGPT to recreate the human-written text several times and selected the best response. For robustness, we used two sets of human and AI-written paragraphs (labeled A and B), instead of relying on one. Our analysis includes data from all paragraphs, allowing for differences in average trust based on the paragraph set reviewed.

The following paragraph was written by a human.

Please read it and then answer the questions below. You have up to 3 minutes to read the paragraph and answer the questions.

The English clergyman Joseph Butler, in his very influential Analogy of Religion (1736), called probability "the very guide of life." The phrase, however, did not refer to mathematical calculation but merely to the judgments made where rational demonstration is impossible. The word probability was used in relation to the mathematics of chance in 1662 in the Logic of Port-Royal, written by Pascal's fellow Jansenists, Antoine Arnauld and Pierre Nicole. But from medieval times to the 18th century and even into the 19th, a probable belief was most often merely one that seemed plausible, came on good authority, or was worthy of approval. Probability, in this sense, was emphasized in England and France from the late 17th century as an answer to skepticism.

Fig. 1. Human-Written Paragraph A.

Table 1
Treatments.

Human-written (informed)	Human-written (uninformed)
AI-written (informed)	AI-written (uninformed)

Table 2
Number of Participants.

	Paragraph A	Paragraph B	Total
Human-written (informed)	78	86	164
Human-written (uninformed)	50	39	89
AI-written (informed)	79	69	148
AI-written (uninformed)	43	56	99

We use these answers to create our measure of trust.[5]

Participants earned $0.50 for correctly answering the question about errors but could also pay $0.20 (or 40 % of the total reward) to fact-check the text and receive $0.30 for sure. If a participant does not fact-check, then their expected payment (expected value) is:

$$EV[nofactcheck] = \$0.50 * P(correct).$$

where $P(correct)$ is the participant's subjective probability that their guess is correct. If a participant does fact check, their expected payment is: $EV\ [fact\ check] = \$0.30$.

$$EV[nofactcheck] > EV[factcheck]$$

when $P(correct) > 0.6$. So, if a risk-neutral participant is at least 60 % certain there are no errors in the text, they maximize their expected earnings by answering "There are not any factual errors in this paragraph" rather than by purchasing a fact-check. Conversely, if a risk-neutral participant is less than 40 % certain there are no errors in the text (e.g. believes that errors likely exist), they maximize their expected earnings by answering "There are factual errors in this paragraph." Risk-neutral participants who are between 40 and 60 % certain there are no errors in the text maximize their expected earnings by purchasing a fact-check. If participants are risk-averse, that might lead to an overall increase in fact-checks, but this is unlikely to differ systematically between treatments.

We categorize responses as follows: no errors is "High trust," errors is "Low trust," and factcheck requests are labeled "Medium trust."

The experiment concluded with a demographic survey and questions about familiarity with ChatGPT.

In June-July 2023, we recruited 500 participants via Prolific.[6] Table 2 shows the number of participants in each treatment.

So as not to rely exclusively on a single paragraph, we used two human-written paragraphs and derived two AI-written paragraphs from them. Each of the paragraphs included details about the history of statistics. We only used paragraphs that did not include controversial facts or polarizing topics. None of the paragraphs contain factual errors, and we did not give participants information about the likelihood that their paragraph had factual errors. In our analysis, we allow average trust to vary depending on which passage the participant viewed. Each participant only saw one paragraph.

4. Hypotheses

Hypothesis 1. There will be higher trust in passages believed to have been written by a human author as opposed to generated by AI.

In the Informed treatments, this would result in higher trust among participants who are informed about human authorship than among those who are informed about AI-authorship. In the Uninformed treatments, this would result in higher trust among participants who believe the passage is human-written than among those who believe the passage is AI-written. We expected this because LLMs are a relatively new technology, and participants might be familiar with published writing by humans being generally reliable. Trust, in this case, measures confidence in objective factual accuracy.

Hypothesis 2. Among those who are informed that the text is written by AI, trust levels will decrease with the age of the participant.

We collect information on age to test whether older participants are less trusting of this new technology.

Hypothesis 3. Among those who are informed that the text is written by AI, those who have used AI for writing will trust AI-written text less than those who have not used AI for writing.

Participants who have used ChatGPT for writing might have directly

[5] Our approach is similar to the ternary list method proposed by Healy and Leo (2024).

[6] Participants in our sample are at least 18 years old, live in the United States, and have at least a 75 % approval rating on Prolific. In addition to a base payment of $1 for completing the survey, we paid performance bonuses to participants who completed the questions. The average participant earned a total of $1.35 and spent 4 minutes and 44 seconds completing the survey ($17.12/hour). As an attention check, participants had to correctly count the number of sentences in the paragraph. We pre-registered the experiment in the American Economic Association's RCT registry.

observed it making a mistake. Therefore, we test whether there is a significant effect on trust in this context.

5. Results

The survey experiment, which we conducted online in June-July 2023, included American adults registered with Prolific. Table 3 contains descriptive statistics of the sample. More than 80 % of participants had heard of ChatGPT but only 35 % had used it to write.

We show the balance of our samples between treatments in Table A1.

We instructed ChatGPT to mimic the tone and language of the human-written paragraphs as closely as possible. We did not expect that uninformed participants will be more or less trusting simply based on the style of writing (consistent with Casal and Kessler (2023)). We only expected to see a difference based on perceived authorship.

If the paragraphs written by ChatGPT are of obviously different quality than the human-written paragraphs, then we would not directly be testing trust in the author's ability to not make factual errors. Among the uninformed, there is no difference in trust between those who are shown human-written text and AI-generated text.[7] We test this in column (2) of Table 4, where we restrict the sample to the uninformed and do not find a significant relationship between trust in the paragraphs and whether the paragraphs were actually human or AI-written. This relationship remains insignificant if we treat our measure of trust as a nominal variable and conduct a chi-square test ($\chi^2(2) = 0.6072$, $p = 0.738$).

Result 1a. Participants are not more trusting of human-written text, when informed about authorship.

Fig. 2a shows that the distribution of choices is similar between human and AI authors when participants are explicitly informed of authorship. The modal response is to purchase a fact-check, and participants do not appear significantly more willing to trust a human author than an AI author. In both of the informed treatments, only about 20% of participants have high trust in the accuracy of the paragraph. Being informed of authorship and being asked an artifactual question may have raised suspicion, equally so for AI and human authored paragraphs.

We use ordered logistic regressions to test for differences in trust between treatments.[8] The regression in column (1) of Table 4 confirms our visual analysis of behavior when participants are explicitly informed about authorship. The coefficient on human-written is positive but small and insignificant. This relationship remains insignificant if we treat our measure of trust as a nominal variable and conduct a chi-square test ($\chi^2(2) = 1.9307$, $p = 0.381$). Therefore, we find that, among the informed, trust does vary significantly by authorship.

Result 1b. When participants are not explicitly informed of authorship, there is higher trust in paragraphs believed to have been written by humans.

We find partial support for our first hypothesis when considering only the trust levels of people who were uninformed. Uninformed participants who believe the paragraph they review is human-written are more confident in the accuracy of the paragraph than subjects who believe the paragraph is AI-written. We show this visually in Fig. 2b, where "High trust" is higher for those who believe the paragraph is human-written than for those who believe the paragraph is AI-written. Thus, participants are more trusting of what they believe to be human authors, or they associate trustworthy writing with human authors. However, this relationship is only evident among the uninformed. The coefficient on "Believed human authorship" in column (3) of Table 4 is positive and significant. Therefore, we find that, among the uninformed, those who believe the text is human-written trust the text's factual accuracy significantly more than those who believe the text is AI-written. This relationship remains significant if we treat our measure of trust as a nominal variable and conduct a chi-square test ($\chi^2(2) = 14.308$, $p = 0.001$).

Result 1c. Participants do more fact checking when they are informed about authorship, regardless of whether the author is human or AI.

We did not anticipate the overall difference between the two treatments that we report in Result 1c. Informed participants choose more fact-checks, regardless of whether they are told the author is human or AI. Fig. 2c visually demonstrates that the rate of fact-checking (although not overall trust) is significantly higher in the Informed treatments than in the Uninformed treatments ($\chi^2(1) = 7.33$, $p = 0.007$). This implies that the level of confidence in participants' own abilities to judge factual accuracy is lower when they are explicitly informed about authorship.

We asked several demographic questions to test hypotheses about

Table 3
Sample Description.

	mean	s.d.	min	max
Age	36.77	12.40	18	74
Female	0.48		0	1
4-year college degree or higher	0.56		0	1
Has heard news stories about AI	0.85		0	1
Has seen examples of writing by ChatGPT	0.86		0	1
Has used ChatGPT for writing	0.35		0	1
Observations	500			

Notes: "Age" is measured in years. All other variables in this table are indicator variables, and the means represent the proportions of our sample for which the statements are true.

Table 4
Trust compared between treatments and by authorship belief.

	(1)	(2)	(3)	(4)
Human-written	0.329	0.149	0.0609	0.294
	(0.223)	(0.276)	(0.279)	(0.214)
Uninformed				−0.141
				(0.252)
Human-written × Uninformed				−0.132
				(0.359)
Believed human authorship			1.105***	
			(0.301)	
Paragraph B	−0.330	−0.155	−0.289	−0.256
	(0.222)	(0.276)	(0.281)	(0.172)
Observations	312	188	188	500

Standard errors in parentheses.
*** $p < .01$.
Notes: The dependent variable in each of these ordered logistic regressions is our measure of Trust. Column (1) includes only participants in the Informed treatments, columns (2)-(3) include only participants in the Uninformed treatments, and column (4) includes all participants. "Human-written" is a binary variable that =1 if a participant reviewed human-written text. "Uninformed" is a binary variable that =1 if a participant was not informed about the authorship of the text they reviewed. "Believed human authorship" is a binary variable that =1 if a participant (in the Uninformed treatments) guessed the text they reviewed was human-written. "Paragraph B" is a binary variable that =1 if a participant saw the second set of paragraphs we included for robustness.

[7] A majority (64 %) of participants in these uninformed treatments guessed AI authorship. Although trust does not vary by authorship, 69 % of those in the AI (uninformed) treatment correctly guessed authorship compared to 40 % of those in the Human (uninformed) treatment. This difference is statistically significant (two-tailed t-test; $t = -4.03$, $p < 0.001$). The results indicate that the writing by the AI was not of obviously different quality than human writing.

[8] In this regression model, the categories of the dependent variable are underpinned by a continuous latent variable. In this case, the underlying trust (which can be between 0 (no trust) to 1 (complete trust)) in the factual accuracy of text is the latent variable. While the coefficients themselves are not easily interpreted, we primarily focus on using this regression framework to test for significant differences in trust between treatments.

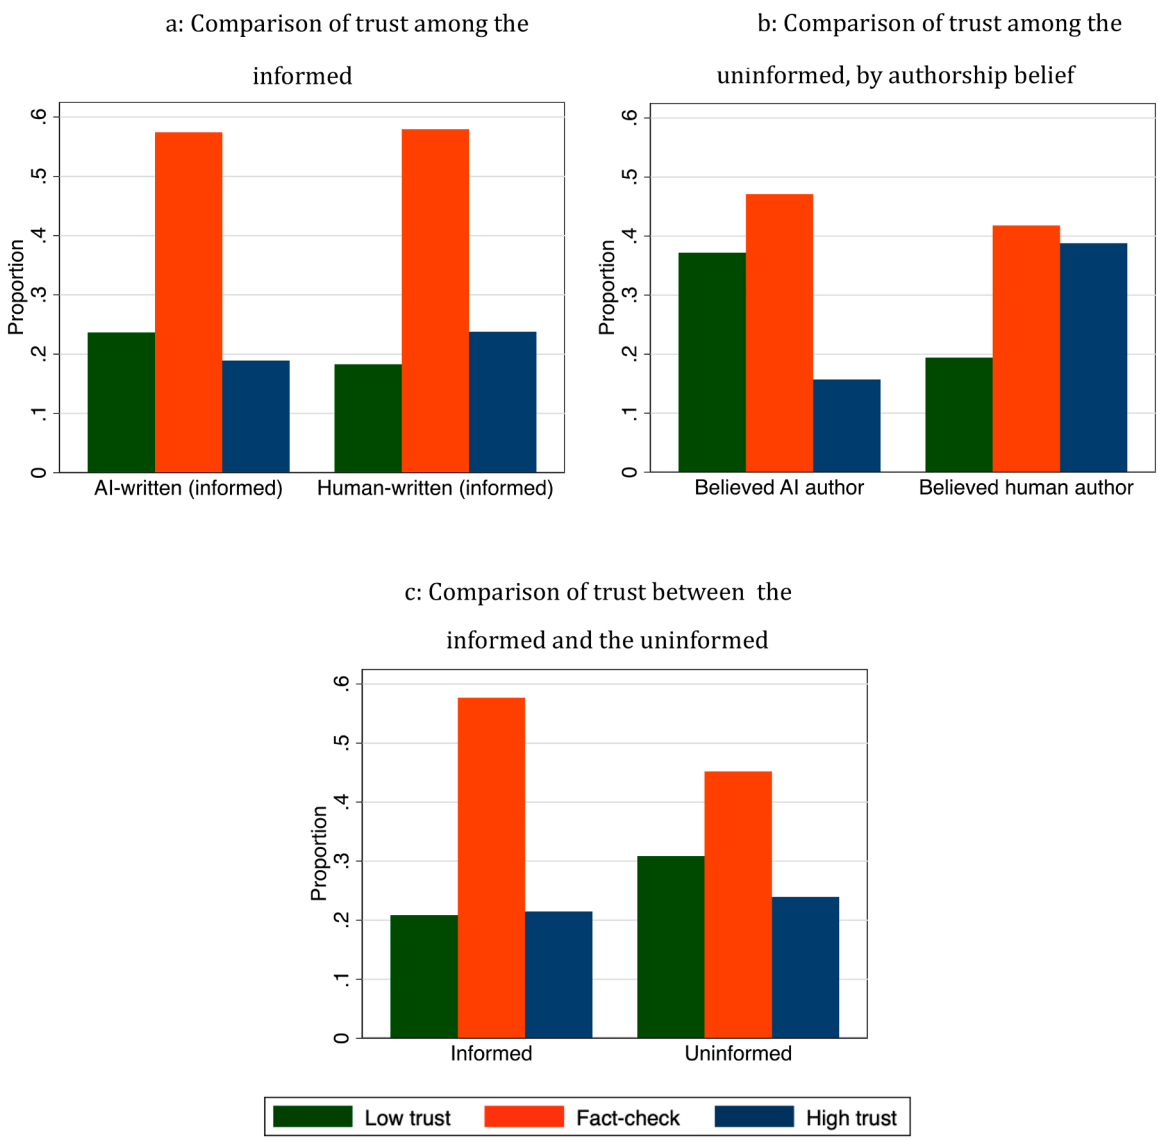

Fig. 2. a: Comparison of trust among the informed; b: Comparison of trust among the uninformed, by authorship belief; c: Comparison of trust between the informed and the uninformed.
Notes: "Informed" is an indicator for treatments where participants are informed of authorship (either human or AI). "Uninformed" is an indicator for treatments where participants are not informed of authorship. "Believed AI author" and "Believed human author" are indicators for whether uninformed participants believe the text they review is written by AI or written by a human. "Low trust" indicates participants believe the text has factual errors, and do not purchase a fact-check. "Fact-check" indicates participants choose to purchase a fact-check. "High trust" indicates participants believe the text does not have factual errors, and do not purchase a fact-check.

the interaction of trust and life experience with technology. The lack of significant results in Table A2 indicates that we cannot reject the null for hypothesis 2 or 3.

Result 2. Trust does not appear to decrease with age in our sample.

Result 3. Prior usage of ChatGPT does not significantly impact trust.

6. Conclusion

Considering that generative AI can produce writing that sounds correct but contains factual errors, it is important to understand how human readers interact with tools like ChatGPT. Our paper contributes to the literature on disinformation and the behavioral aspects of advances in artificial intelligence. We find that when participants believe text was written by a human, they display higher levels of trust in one information condition. However, trust is context-dependent. Participants who we informed about authorship displayed equally low levels of trust toward both human and AI writing.

AI is a powerful tool that will be widely adopted. Policymakers and researchers have expressed concerns about "alignment" in this new era. Korinek (2023) lists the following as one of the "urgent" research questions for economists: "Finally, but perhaps most importantly, how can we best address the AI alignment problem, i.e., ensure that ever-more advanced and potentially super-intelligent AI systems pursue objectives that are aligned with human objectives?" Messages generated by AI might not always be true or convey the most helpful content to align with the goals of humans who query a tool like ChatGPT or who read an AI-generated website. Therefore, it is important to understand how context influences human readers and how to ensure that humans read critically when they encounter AI-generated content.

People tend to be overconfident in their own ability to assess text (Serra-Garcia & Gneezy, 2021). If certain framing prompts people to fact-check, then that could be a promising policy avenue. We find that

people are more likely to fact-check when we explicitly informed them about authorship. Similarly, Pennycook et al. (2021) found that people are more hesitant to share stories on social media if they are prompted to consider accuracy. The European Union has proposed, as one part of a policy to regulate artificial intelligence, a requirement that generative AI comply with transparency requirements such as "Disclosing that the content was generated by AI" (EU AI Act 2024).

Based on our results, certain ways of displaying AI-generated writing may encourage human readers to be alert and fact-check when necessary. This would allow humans to receive the benefits of AI while mitigating the risks of disinformation or bias. Although we did not find significant correlations between the demographic information that we collected and trust, future research can explore whether certain groups are more trusting of AI than others and are therefore more vulnerable to misinformation. We consider this to be an important new field of behavioral research because AI-based tools are quickly changing the world around us but are not infallible.

CRediT authorship contribution statement

Joy Buchanan: Writing – review & editing, Writing – original draft, Visualization, Validation, Supervision, Software, Resources, Project administration, Methodology, Investigation, Funding acquisition, Formal analysis, Data curation, Conceptualization. **William Hickman:** Writing – review & editing, Writing – original draft, Visualization, Validation, Supervision, Software, Resources, Project administration, Methodology, Investigation, Formal analysis, Data curation, Conceptualization.

Data availability

Data will be made available on request.

Supplementary materials

Supplementary material associated with this article can be found, in the online version, at doi:10.1016/j.socec.2024.102239.

References

Athey, S., Cersosimo, M., Koutout, K., and Li, Z. (2023).Emotion-versus reasoning-based drivers of misinformation sharing: A field experiment using text message courses in Kenya. Stanford University Graduate School of Business Research Paper No. 4489759, Available at SSRN: https://ssrn.com/abstract=4489759.

Babin, J.J., and Chauhan, H. (2023). "Chatbot or humanaut? How the source of advice impacts behavior in one-shot social dilemmas". *Working Paper*.

Brynjolfsson, E., Li, D., and Raymond, L.R. (2023). Generative AI at work. National Bureau of Economic Research Working Paper 31161. http://www.nber.org/papers/w31161.

Buchanan, J., Hill, S., & Shapoval, O. (2024). ChatGPT hallucinates non-existent citations: Evidence from economics. *The American Economist, 69*(1), 80–87. https://doi.org/10.1177/05694345231218454

Buchanan, J. (2022). Willingness to be paid: Who trains for tech jobs? *Labour Economics, 79*, Article 102267.

Casal, J. E., & Kessler, M. (2023). Can linguists distinguish between ChatGPT/AI and human writing?: A study of research ethics and academic publishing. *Research Methods in Applied Linguistics, 2*(3), Article 100068. https://doi.org/10.1016/j.rmal.2023.100068

Castelo, N., Bos, M. W., & Lehmann, D. R. (2019). Task-dependent algorithm aversion. *Journal of Marketing Research, 56*(5), 809–825.

Chen, J., & Houser, D. (2017). Promises and lies: Can observers detect deception in written messages. *Experimental Economics, 20*, 396–419.

Chugunova, M., & Sele, D. (2022). We and It: An interdisciplinary review of the experimental evidence on how humans interact with machines. *Journal of Behavioral and Experimental Economics, 99*, Article 101897. https://doi.org/10.1016/j.socec.2022.101897

Epstein, Z., Fang, M. C., Arechar, A. A., & Rand, D. G.. What label should be applied to content produced by generative AI?. https://doi.org/10.31234/osf.io/v4mfz.

EU AI Act. (2024). *EU ai act: first regulation on artificial intelligence — topics — european parliament*. url https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence. visited on 02/15/2024.

Fumagalli, E., Rezaei, S., & Salomons, A. (2022). OK computer: Worker perceptions of algorithmic recruitment. *Research Policy, 51*(2), Article 104420.

Gillespie, N., Lockey, S., Curtis, C., Pool, J., and Akbari, A. (2023). Trust in artificial intelligence: a global study. The University of Queensland and KPMG Australia. doi:10.14264/00d3c94.

Healy, Paul J. and Leo, Greg, Minimal Experiments. Available at SSRN: https://ssrn.com/abstract=4730244.

Ji, Z., Frieske, R., Yu, T., Su, D., Xu, Y., … Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys, 55*(12), 1–38.

Köbis, N., & Mossink, L. D. (2021). Artificial intelligence versus Maya Angelou: Experimental evidence that people cannot differentiate AI-generated from human-written poetry. *Computers in human behavior, 114*, Article 106553.

Korinek, A. (2023). Generative AI for economic research: use cases and implications for economists. *Journal of Economic Literature, 61*(4), 1281–1317. url https://www.aeaweb.org/articles?id=10.1257/jel.20231736.

Leib, M., Köbis, N., Rilke, R. M., Hagens, M., & Irlenbusch, B. (2023). Corrupted by algorithms? How ai-generated and human-written advice shape (Dis)honesty. *The Economic Journal, 134*(658), 766–784. https://doi.org/10.1093/ej/uead056

March, C. (2021). Strategic interactions between humans and artificial intelligence: Lessons from experiments with computer players. *Journal of Economic Psychology, 87*, Article 102426. https://doi.org/10.1016/j.joep.2021.102426

OpenAI. (2023). GPT-4V(ision) system card. https://openai.com/index/gpt-4v-system-card/.

Pennycook, G., Epstein, Z., Mosleh, M., Arechar, A. A., Eckles, D., & Rand, D. G. (2021). Shifting attention to accuracy can reduce misinformation online. *Nature, 592*(7855), 590–595.

Porter, T. M. (2023). Probability and statistics. Encyclopedia Britannica, https://www.britannica.com/science/probability. Accessed 19 May 2023.

Serra-Garcia, M., & Gneezy, U. (2021). Mistakes, overconfidence, and the effect of sharing on detecting lies. *American Economic Review, 111*(10), 3160–3183.

Smith, V. L., & Wilson, B. J. (2019). *Humanomics: moral sentiments and the wealth of nations for the twenty-first century*. Cambridge, UK: Cambridge University Press.

Smith, V. L. (1994). Economics in the Laboratory. *Journal of economic perspectives, 8*(1), 113–131.

Spitale, G., Biller-Andorno, N., & Germani, F. (2023). AI model GPT-3 (dis) informs us better than humans. *Science Advances, 9*(26), eadh1850.

Sunstein, C. R., & Reisch, L.. Do people like algorithms? a research strategy. https://ssrn.com/abstract=4544749.

Zhang, M., Press, O., Merrill, W., Liu, Al, and Smith, N.A. (2023). How language model hallucinations can snowball. https://arxiv.org/abs/2305.13534.